Qwen3-32 模型服务性能测试结果（Agent场景）

Model Service Benchmark Report

Test Config

- Script version: `5`
- Generated at: `2026-03-22T11:01:51`
- Base URL: `http://127.0.0.1:8000/v1\`
- Mode: `bench`
- Models: `Qwen3-32B`
- Context sizes: `128,512,2048,4096,8192`
- Bench testset: `model_service_testset.json`
- Thinking mode: `no_think`
- Max output tokens: `128`
- Temperature: `0.0`
- Timeout per request: `1800s`
- First token timeout: `60s`

Tokenizer Status

- `Qwen3-32B` -> `rough_estimate (No module named 'transformers')`

Environment Baseline

- Hostname: `hw-xn-cd-rbt-4-dvznt-001`
- Platform: `linux`
- OS: `Ubuntu 20.04.6 LTS`
- Kernel: `Linux hw-xn-cd-rbt-4-dvznt-001 5.4.0-162-generic #179-Ubuntu SMP Mon Aug 14 08:51:31 UTC 2023 x86_64 x86_64 x86_64 GNU/Linux`
- Python: `3.9.5`

- CPU model: `Intel Xeon Processor (Cascadelake)`
- CPU logical cores: `32`
- CPU physical cores: `16`
- Total memory (GB): `251.889`
- Root disk total/free (GB): `245.956` / `73.609`
- GPU summary: `GPU 0: Tesla V100S-PCIE-32GB (UUID: GPU-3beb1d35-b4eb-59a5-89d8-c314944a6847); GPU 1: Tesla V100S-PCIE-32GB (UUID: GPU-dfefd7fe-ee7c-0a25-2da2-03666ec762bb); GPU 2: Tesla V100S-PCIE-32GB (UUID: GPU-d7586c8e-d0ce-6736-491e-625344696468); GPU 3: Tesla V100S-PCIE-32GB (UUID: GPU-086b7df6-bf85-67f1-13f0-6e92592ad511)`
- nvidia-smi summary: `Tesla V100S-PCIE-32GB, 32768, 535.230.02; Tesla V100S-PCIE-32GB, 32768, 535.230.02; Tesla V100S-PCIE-32GB, 32768, 535.230.02; Tesla V100S-PCIE-32GB, 32768, 535.230.02`

Model Summary

Model	Smoke Success	Bench Success	Avg TTFT (s)	Max TTFT (s)	Avg Output TPS	Avg GPU Util (%)	Avg KV Cache (%)
Qwen3-32B	-	16/16	11.849	15.649	17.292	92.859	0.000

Bench Results

Model	Case	Prompt Tokens (think/no_think)	Think OK	NoThink OK	Think TTFT (s)	NoThink TTFT (s)	Think Total (s)	NoThink Total (s)	Think Completion	NoThink Completion	NoThink Output

Qwen3-32B	step_3	7790/7791	Y	Y	9.924	9.919	18.976	10.277	208	4	OK。
Qwen3-32B	step_5	11507/11509	Y	Y	10.064	9.913	19.164	10.198	200	4	OK
Qwen3-32B	step_8	16739/16740	Y	Y	15.578	15.649	24.642	15.933	195	4	OK
Qwen3-32B	step_10	19385/19386	Y	Y	11.132	11.158	20.197	11.440	181	4	OK
Qwen3-32B	step_12	21650/21651	Y	Y	10.013	10.017	19.067	10.297	190	4	OK
Qwen3-32B	step_15	25037/25038	Y	Y	15.395	15.435	24.416	15.714	182	4	OK
Qwen3-32B	step_18	27892/27894	Y	Y	14.492	14.497	23.517	14.782	177	4	OK
Qwen3-32B	step_20	29425/29427	Y	Y	8.403	7.992	17.443	8.285	177	4	OK

Bench Per-Mode Details

Model	Case	Mode	Prompt Tokens	Success	HTTP	TTFT (s)	Total (s)	Decode (s)	Completion Tokens	Output TPS	Chunks	Output Preview	Service Metrics
Qwen3-32B	step_3	think	7790	Y	200	9.924	18.976	9.053	208	22.976	129	好的，我现在需要处理用户的问题，关于检索链路在大促压测期间出现	kv_cache=0.00% (before=0.00%) prefix_cache_hit_rate=0.00% gpu_util_avg=86.00% gpu_mem_uti

												的延迟抖动和局部召回空结果的情况。用户已经提供了详细的运行手册检索结果、最近	l_avg=86.89% gpu_mem_used_total_mb=113888

的变更记录以及监控摘要。我的任务是根据这些信息继续调查问题，并按照ReAct运维...

Qwen3-32B	step_3	no_think	7791	Y	200	9.919	10.277	0.358	4	11.168	8	OK。	kv_cache=0.00% (before=0.00%) prefix_cache_hit_rate=0.62% gpu_util_avg=88.00% gpu_mem_util_avg=86.89% gpu_mem

													_used_total_mb=113888
Qwen3-32B	step_5	think	11507	Y	200	10.064	19.164	9.100	200	21.979	129	好的，我现在需要处理用户关于检索链路在大促压测期间出现延迟抖动	kv_cache=0.00% (before=0.00%) prefix_cache_hit_rate=64.51% gpu_util_avg=84.50% gpu_mem_uti

												和局部召回空结果的问题。根据之前的ReAct运维代理的执行历史，已经完成了运行手册	l_avg=86.89% gpu_mem_used_total_mb=113888

检索、最近变更检查、监控摘要分析、关键错误日志查看以及配置差异对比。现在需...

Qwen3-32B	step_5	no_think	11509	Y	200	9.913	10.198	0.285	4	14.042	7	OK	kv_cache=0.00% (before=0.00%) prefix_cache_hit_rate=64.49% gpu_util_avg=95.25% gpu_mem_util_avg=86.89% gpu_mem

													_used_total_mb=113888
Qwen3-32B	step_8	think	16739	Y	200	15.578	24.642	9.064	195	21.514	129	好的，我现在需要处理用户关于检索链路在大促压测期间出现延迟抖动	kv_cache=0.00% (before=0.00%) prefix_cache_hit_rate=70.78% gpu_util_avg=91.50% gpu_mem_uti

												和局部召回空结果的问题。根据之前的对话历史，用户已经进行了多个步骤的排查，包	l_avg=86.89% gpu_mem_used_total_mb=113888

括运行手册检索、最近变更检查、监控摘要分析、错误日志查看、配置差异对比、工单

												反馈...	
Qwen3-32B	step_8	no_think	16740	Y	200	15.649	15.933	0.284	4	14.106	7	OK	kv_cache=0.00% (before=0.00%) prefix_cache_hit_rate=70.76% gpu_util_avg=94.50% gpu_mem_util_avg=86.89%

													gpu_mem_used_total_mb=113888
Qwen3-32B	step_10	think	19385	Y	200	11.132	20.197	9.065	181	19.968	129	好的，我现在需要处理用户的问题，即检索链路在大促压测期间	kv_cache=0.00% (before=0.00%) prefix_cache_hit_rate=85.45% gpu_util_avg=92.75%

												出现延迟抖动和局部召回空结果的情况。用户希望我像一个ReAct运维代理一样逐步调查	gpu_mem_util_avg=86.89% gpu_mem_used_total_mb=113888

，优先检查运行手册、最近变更、监控、日志、配置、工单和集群事件，再决定下一步

												° ...	
Qwen3-32B	step_10	no_think	19386	Y	200	11.158	11.440	0.283	4	14.151	7	OK	kv_cache=0.00% (before=0.00%) prefix_cache_hit_rate=85.44% gpu_util_avg=94.25% gpu_mem_util_avg=86.89%

													gpu_mem_used_total_mb=113888
Qwen3-32B	step_12	think	21650	Y	200	10.013	19.067	9.054	190	20.985	129	好的，我现在需要处理用户的问题，即检索链路在大促压测期间	kv_cache=0.00% (before=0.00%) prefix_cache_hit_rate=89.84% gpu_util_avg=93.50%

												出现延迟抖动和局部召回空结果的情况。根据之前的工具调用和结果，我需要综合所有	gpu_mem_util_avg=86.89% gpu_mem_used_total_mb=113888

信息来确定根本原因并提出解决方案。首先，回顾之前的步骤，用户已经通过多个工具

												调用来收...	
Qwen3-32B	step_12	no_think	21651	Y	200	10.017	10.297	0.280	4	14.306	7	OK	kv_cache=0.00% (before=0.00%) prefix_cache_hit_rate=89.82% gpu_util_avg=100.00% gpu_mem_util_av

													g=86.89% gpu_mem_used_total_mb=113888
Qwen3-32B	step_15	think	25037	Y	200	15.395	24.416	9.021	182	20.175	129	好的，我现在需要处理用户的问题，即检索链路在大促	kv_cache=0.00% (before=0.00%) prefix_cache_hit_rate=87.32% gpu_util_avg=9

												压测期间出现延迟抖动和局部召回空结果的情况。根据之前的工具调用和结果，我需要	5.00% gpu_mem_util_avg=86.89% gpu_mem_used_total_mb=113888

综合所有信息来确定根本原因并提出解决方案。首先，回顾之前的步骤，用户已经通过

												多个工具调用来收...	
Qwen3-32B	step_15	no_think	25038	Y	200	15.435	15.714	0.278	4	14.379	7	OK	kv_cache=0.00% (before=0.00%) prefix_cache_hit_rate=87.30% gpu_util_avg=99.25% gpu_m

													em_util_avg=86.89% gpu_mem_used_total_mb=113888
Qwen3-32B	step_18	think	27892	Y	200	14.492	23.517	9.025	177	19.612	129	好的，我现在需要处理用户的问题，即检索链路	kv_cache=0.00% (before=0.00%) prefix_cache_hit_rate=90.28%

												在大促压测期间出现延迟抖动和局部召回空结果的情况。用户希望我像一个ReAct运维代	gpu_util_avg=80.75% gpu_mem_util_avg=86.89% gpu_mem_used_total_mb=113888

理一样逐步调查，优先检查运行手册、最近变更、监控、日志、配置、工单和集群事件

												，再决定下一步。...	
Qwen3-32B	step_18	no_think	27894	Y	200	14.497	14.782	0.284	4	14.062	7	OK	kv_cache=0.00% (before=0.00%) prefix_cache_hit_rate=90.26% gpu_util_avg=100.00%

													gpu_mem_util_avg=86.89% gpu_mem_used_total_mb=113888
Qwen3-32B	step_20	think	29425	Y	200	8.403	17.443	9.040	177	19.580	129	好的，我现在需要处理用户的问题，即检索	kv_cache=0.00% (before=0.00%) prefix_cache_hit_rate=95

												链路在大促压测期间出现延迟抖动和局部召回空结果的情况。用户希望我像一个ReAct运	.06%gpu_util_avg=90.50%gpu_mem_util_avg=86.89%gpu_mem_used_total_mb=113888

维代理一样逐步调查，优先检查运行手册、最近变更、监控、日志、配置、工单和集群

												事件，再决定下一步。...	
Qwen3-32B	step_20	no_think	29427	Y	200	7.992	8.285	0.293	4	13.670	7	OK	kv_cache=0.00% (before=0.00%) prefix_cache_hit_rate=95.10% gpu_util_avg=100.

													00% gpu_mem_util_avg=86.89% gpu_mem_used_total_mb=113888

Raw Results

复制代码

```
{
  "base_url": "http://127.0.0.1:8000/v1",
  "mode": "bench",
  "models": [
    {
      "name": "Qwen3-32B",
      "tokenizer": "/data/models/Qwen3-32B"
    }
  ],
  "context_sizes": [
    128,
    512,
    2048,
```

```
    4096,
    8192
  ],
  "testset_path": "model_service_testset.json",
  "thinking_mode": "no_think",
  "max_tokens": 128,
  "temperature": 0.0,
  "timeout": 1800,
  "first_token_timeout": 60,
  "environment": {
    "hostname": "hw-xn-cd-rbt-4-dvznt-001",
    "platform": "linux",
    "kernel": "Linux hw-xn-cd-rbt-4-dvznt-001 5.4.0-162-generic #179-
Ubuntu SMP Mon Aug 14 08:51:31 UTC 2023 x86_64 x86_64 x86_64 GNU/Linux",
    "os_release": "Ubuntu 20.04.6 LTS",
    "python_version": "3.9.5",
    "cpu_model": "Intel Xeon Processor (Cascadelake)",
    "cpu_cores_logical": 32,
    "cpu_cores_physical": 16,
    "total_memory_gb": 251.88890838623047,
    "disk_root_total_gb": 245.95623016357422,
    "disk_root_free_gb": 73.60884475708008,
    "gpu_summary": [
      "GPU 0: Tesla V100S-PCIE-32GB (UUID: GPU-3beb1d35-b4eb-59a5-89d8-
c314944a6847)",
      "GPU 1: Tesla V100S-PCIE-32GB (UUID: GPU-dfefd7fe-ee7c-0a25-2da2-
03666ec762bb)",
      "GPU 2: Tesla V100S-PCIE-32GB (UUID: GPU-d7586c8e-d0ce-6736-491e-
625344696468)",
      "GPU 3: Tesla V100S-PCIE-32GB (UUID: GPU-086b7df6-bf85-67f1-13f0-
6e92592ad511)"
    ],
    "nvidia_smi_summary": "Tesla V100S-PCIE-32GB, 32768, 535.230.02;
Tesla V100S-PCIE-32GB, 32768, 535.230.02; Tesla V100S-PCIE-32GB, 32768,
535.230.02; Tesla V100S-PCIE-32GB, 32768, 535.230.02",

    "raw_notes": [ ]

  },

  "smoke_results": [ ],
```

```
  "bench_results": [
    {
      "model": "Qwen3-32B",
      "case_name": "step_3",
      "thinking_mode": "think",
      "context_tokens_target": 7790,
      "success": true,
      "http_status": 200,
      "error": "",
      "ttft_seconds": 9.92360903462395,
      "total_seconds": 18.976425798609853,
      "decode_seconds": 9.052816763985902,
      "raw_chunks": 129,
      "prompt_tokens_estimated": 7790,
      "completion_tokens_estimated": 208,
      "prompt_tokens_reported": null,
      "completion_tokens_reported": null,
      "total_tokens_reported": null,
      "output_tps_estimated": 22.976274172196856,
      "output_preview": "<think> 好的，我现在需要处理用户的问题，关于检索链路在大促压测期间出现的延迟抖动和局部召回空结果的情况。用户已经提供了详细的运行手册检索结果、最近的变更记录以及监控摘要。我的任务是根据这些信息继续调查问题，并按照ReAct运维...",
      "metrics_before": {
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
        "prefix_cache_queries": 6609.0,
        "prefix_cache_hits": 0.0,
        "gpu_utilization_avg": 0.0,
        "gpu_memory_used_mb_total": 113528.0,
        "gpu_memory_utilization_avg": 86.614990234375,

        "raw_notes": [ ]

      },
      "metrics_after": {
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
```

```
      "prefix_cache_queries": 14334.0,
      "prefix_cache_hits": 0.0,
      "gpu_utilization_avg": 86.0,
      "gpu_memory_used_mb_total": 113888.0,
      "gpu_memory_utilization_avg": 86.8896484375,

      "raw_notes": [ ]

    }
  },
  {
    "model": "Qwen3-32B",
    "case_name": "step_3",
    "thinking_mode": "no_think",
    "context_tokens_target": 7791,
    "success": true,
    "http_status": 200,
    "error": "",
    "ttft_seconds": 9.919018597807735,
    "total_seconds": 10.277193604037166,
    "decode_seconds": 0.35817500622943044,
    "raw_chunks": 8,
    "prompt_tokens_estimated": 7791,
    "completion_tokens_estimated": 4,
    "prompt_tokens_reported": null,
    "completion_tokens_reported": null,
    "total_tokens_reported": null,
    "output_tps_estimated": 11.167725079727601,
    "output_preview": "<think> </think> OK。",
    "metrics_before": {
      "collected": true,
      "source": "vllm_metrics+nvidia_smi",
      "kv_cache_usage_perc": 0.0,
      "prefix_cache_queries": 14334.0,
      "prefix_cache_hits": 0.0,
      "gpu_utilization_avg": 86.0,
      "gpu_memory_used_mb_total": 113888.0,
      "gpu_memory_utilization_avg": 86.8896484375,

      "raw_notes": [ ]
```

```
      },
      "metrics_after": {
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
        "prefix_cache_queries": 22063.0,
        "prefix_cache_hits": 48.0,
        "gpu_utilization_avg": 88.0,
        "gpu_memory_used_mb_total": 113888.0,
        "gpu_memory_utilization_avg": 86.8896484375,

        "raw_notes": [ ]

      }
    },
    {
      "model": "Qwen3-32B",
      "case_name": "step_5",
      "thinking_mode": "think",
      "context_tokens_target": 11507,
      "success": true,
      "http_status": 200,
      "error": "",
      "ttft_seconds": 10.06434306083247,
      "total_seconds": 19.163893630728126,
      "decode_seconds": 9.099550569895655,
      "raw_chunks": 129,
      "prompt_tokens_estimated": 11507,
      "completion_tokens_estimated": 200,
      "prompt_tokens_reported": null,
      "completion_tokens_reported": null,
      "total_tokens_reported": null,
      "output_tps_estimated": 21.97910748050202,
      "output_preview": "<think> 好的，我现在需要处理用户关于检索链路在大促压测期间出现延迟抖动和局部召回空结果的问题。根据之前的ReAct运维代理的执行历史，已经完成了运行手册检索、最近变更检查、监控摘要分析、关键错误日志查看以及配置差异对比。现在需...",
      "metrics_before": {
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
```

```
      "prefix_cache_queries": 22063.0,
      "prefix_cache_hits": 48.0,
      "gpu_utilization_avg": 88.0,
      "gpu_memory_used_mb_total": 113888.0,
      "gpu_memory_utilization_avg": 86.8896484375,

      "raw_notes": [ ]

    },
    "metrics_after": {
      "collected": true,
      "source": "vllm_metrics+nvidia_smi",
      "kv_cache_usage_perc": 0.0,
      "prefix_cache_queries": 34018.0,
      "prefix_cache_hits": 7760.0,
      "gpu_utilization_avg": 84.5,
      "gpu_memory_used_mb_total": 113888.0,
      "gpu_memory_utilization_avg": 86.8896484375,

      "raw_notes": [ ]

    }
  },
  {
    "model": "Qwen3-32B",
    "case_name": "step_5",
    "thinking_mode": "no_think",
    "context_tokens_target": 11509,
    "success": true,
    "http_status": 200,
    "error": "",
    "ttft_seconds": 9.912894133944064,
    "total_seconds": 10.197755597997457,
    "decode_seconds": 0.2848614640533924,
    "raw_chunks": 7,
    "prompt_tokens_estimated": 11509,
    "completion_tokens_estimated": 4,
    "prompt_tokens_reported": null,
    "completion_tokens_reported": null,
    "total_tokens_reported": null,
    "output_tps_estimated": 14.04191336758091,
```

```
    "output_preview": "<think> </think> OK",
    "metrics_before": {
      "collected": true,
      "source": "vllm_metrics+nvidia_smi",
      "kv_cache_usage_perc": 0.0,
      "prefix_cache_queries": 34018.0,
      "prefix_cache_hits": 7760.0,
      "gpu_utilization_avg": 67.0,
      "gpu_memory_used_mb_total": 113888.0,
      "gpu_memory_utilization_avg": 86.8896484375,

      "raw_notes": [ ]

    },
    "metrics_after": {
      "collected": true,
      "source": "vllm_metrics+nvidia_smi",
      "kv_cache_usage_perc": 0.0,
      "prefix_cache_queries": 45977.0,
      "prefix_cache_hits": 15472.0,
      "gpu_utilization_avg": 95.25,
      "gpu_memory_used_mb_total": 113888.0,
      "gpu_memory_utilization_avg": 86.8896484375,

      "raw_notes": [ ]

    }
  },
  {
    "model": "Qwen3-32B",
    "case_name": "step_8",
    "thinking_mode": "think",
    "context_tokens_target": 16739,
    "success": true,
    "http_status": 200,
    "error": "",
    "ttft_seconds": 15.577955306041986,
    "total_seconds": 24.641841178759933,
    "decode_seconds": 9.063885872717947,
    "raw_chunks": 129,
    "prompt_tokens_estimated": 16739,
```

```
      "completion_tokens_estimated": 195,
      "prompt_tokens_reported": null,
      "completion_tokens_reported": null,
      "total_tokens_reported": null,
      "output_tps_estimated": 21.51395138225921,
      "output_preview": "<think> 好的，我现在需要处理用户关于检索链路在大促压测期间出现延迟抖动和局部召回空结果的问题。根据之前的对话历史，用户已经进行了多个步骤的排查，包括运行手册检索、最近变更检查、监控摘要分析、错误日志查看、配置差异对比、工单反馈...",
      "metrics_before": {
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
        "prefix_cache_queries": 45977.0,
        "prefix_cache_hits": 15472.0,
        "gpu_utilization_avg": 89.5,
        "gpu_memory_used_mb_total": 113888.0,
        "gpu_memory_utilization_avg": 86.8896484375,

        "raw_notes": [ ]

      },
      "metrics_after": {
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
        "prefix_cache_queries": 62841.0,
        "prefix_cache_hits": 27408.0,
        "gpu_utilization_avg": 91.5,
        "gpu_memory_used_mb_total": 113888.0,
        "gpu_memory_utilization_avg": 86.8896484375,

        "raw_notes": [ ]

      }
    },
    {
      "model": "Qwen3-32B",
      "case_name": "step_8",
      "thinking_mode": "no_think",
      "context_tokens_target": 16740,
```

```
      "success": true,
      "http_status": 200,
      "error": "",
      "ttft_seconds": 15.64940412901342,
      "total_seconds": 15.932977526914328,
      "decode_seconds": 0.2835733979009092,
      "raw_chunks": 7,
      "prompt_tokens_estimated": 16740,
      "completion_tokens_estimated": 4,
      "prompt_tokens_reported": null,
      "completion_tokens_reported": null,
      "total_tokens_reported": null,
      "output_tps_estimated": 14.10569549051193,
      "output_preview": "<think> </think> OK",
      "metrics_before": {
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
        "prefix_cache_queries": 62841.0,
        "prefix_cache_hits": 27408.0,
        "gpu_utilization_avg": 84.0,
        "gpu_memory_used_mb_total": 113888.0,
        "gpu_memory_utilization_avg": 86.8896484375,

        "raw_notes": [ ]

      },
      "metrics_after": {
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
        "prefix_cache_queries": 79709.0,
        "prefix_cache_hits": 39344.0,
        "gpu_utilization_avg": 94.5,
        "gpu_memory_used_mb_total": 113888.0,
        "gpu_memory_utilization_avg": 86.8896484375,

        "raw_notes": [ ]

      }
    },
```

```
    {
      "model": "Qwen3-32B",
      "case_name": "step_10",
      "thinking_mode": "think",
      "context_tokens_target": 19385,
      "success": true,
      "http_status": 200,
      "error": "",
      "ttft_seconds": 11.132444838061929,
      "total_seconds": 20.19700059387833,
      "decode_seconds": 9.0645557558164,
      "raw_chunks": 129,
      "prompt_tokens_estimated": 19385,
      "completion_tokens_estimated": 181,
      "prompt_tokens_reported": null,
      "completion_tokens_reported": null,
      "total_tokens_reported": null,
      "output_tps_estimated": 19.967884237885436,
      "output_preview": "<think> 好的，我现在需要处理用户的问题，即检索链路在大促压测期间出现延迟抖动和局部召回空结果的情况。用户希望我像一个ReAct运维代理一样逐步调查，优先检查运行手册、最近变更、监控、日志、配置、工单和集群事件，再决定下一步。 ...",
      "metrics_before": {
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
        "prefix_cache_queries": 79709.0,
        "prefix_cache_hits": 39344.0,
        "gpu_utilization_avg": 94.5,
        "gpu_memory_used_mb_total": 113888.0,
        "gpu_memory_utilization_avg": 86.8896484375,

        "raw_notes": [ ]

      },
      "metrics_after": {
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
        "prefix_cache_queries": 99425.0,
        "prefix_cache_hits": 56192.0,
```

```
      "gpu_utilization_avg": 92.75,
      "gpu_memory_used_mb_total": 113888.0,
      "gpu_memory_utilization_avg": 86.8896484375,

      "raw_notes": [ ]

    }
  },
  {
    "model": "Qwen3-32B",
    "case_name": "step_10",
    "thinking_mode": "no_think",
    "context_tokens_target": 19386,
    "success": true,
    "http_status": 200,
    "error": "",
    "ttft_seconds": 11.157778067979962,
    "total_seconds": 11.440452800132334,
    "decode_seconds": 0.2826747321523726,
    "raw_chunks": 7,
    "prompt_tokens_estimated": 19386,
    "completion_tokens_estimated": 4,
    "prompt_tokens_reported": null,
    "completion_tokens_reported": null,
    "total_tokens_reported": null,
    "output_tps_estimated": 14.150539630984229,
    "output_preview": "<think> </think> OK",
    "metrics_before": {
      "collected": true,
      "source": "vllm_metrics+nvidia_smi",
      "kv_cache_usage_perc": 0.0,
      "prefix_cache_queries": 99425.0,
      "prefix_cache_hits": 56192.0,
      "gpu_utilization_avg": 82.0,
      "gpu_memory_used_mb_total": 113888.0,
      "gpu_memory_utilization_avg": 86.8896484375,

      "raw_notes": [ ]

    },
    "metrics_after": {
```

```
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
        "prefix_cache_queries": 119145.0,
        "prefix_cache_hits": 73040.0,
        "gpu_utilization_avg": 94.25,
        "gpu_memory_used_mb_total": 113888.0,
        "gpu_memory_utilization_avg": 86.8896484375,

        "raw_notes": [ ]

      }
    },
    {
      "model": "Qwen3-32B",
      "case_name": "step_12",
      "thinking_mode": "think",
      "context_tokens_target": 21650,
      "success": true,
      "http_status": 200,
      "error": "",
      "ttft_seconds": 10.012921478133649,
      "total_seconds": 19.066865098197013,
      "decode_seconds": 9.053943620063365,
      "raw_chunks": 129,
      "prompt_tokens_estimated": 21650,
      "completion_tokens_estimated": 190,
      "prompt_tokens_reported": null,
      "completion_tokens_reported": null,
      "total_tokens_reported": null,
      "output_tps_estimated": 20.98533058886778,
      "output_preview": "<think> 好的，我现在需要处理用户的问题，即检索链路在大促压测期间出现延迟抖动和局部召回空结果的情况。根据之前的工具调用和结果，我需要综合所有信息来确定根本原因并提出解决方案。 首先，回顾之前的步骤，用户已经通过多个工具调用来收...",
      "metrics_before": {
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
        "prefix_cache_queries": 119145.0,
        "prefix_cache_hits": 73040.0,
```

```
      "gpu_utilization_avg": 93.75,
      "gpu_memory_used_mb_total": 113888.0,
      "gpu_memory_utilization_avg": 86.8896484375,

      "raw_notes": [ ]

    },
    "metrics_after": {
      "collected": true,
      "source": "vllm_metrics+nvidia_smi",
      "kv_cache_usage_perc": 0.0,
      "prefix_cache_queries": 141069.0,
      "prefix_cache_hits": 92736.0,
      "gpu_utilization_avg": 93.5,
      "gpu_memory_used_mb_total": 113888.0,
      "gpu_memory_utilization_avg": 86.8896484375,

      "raw_notes": [ ]

    }
  },
  {
    "model": "Qwen3-32B",
    "case_name": "step_12",
    "thinking_mode": "no_think",
    "context_tokens_target": 21651,
    "success": true,
    "http_status": 200,
    "error": "",
    "ttft_seconds": 10.017023098189384,
    "total_seconds": 10.296626781113446,
    "decode_seconds": 0.279603682924062,
    "raw_chunks": 7,
    "prompt_tokens_estimated": 21651,
    "completion_tokens_estimated": 4,
    "prompt_tokens_reported": null,
    "completion_tokens_reported": null,
    "total_tokens_reported": null,
    "output_tps_estimated": 14.305963205378687,
    "output_preview": "<think> </think> OK",
    "metrics_before": {
```

```
      "collected": true,
      "source": "vllm_metrics+nvidia_smi",
      "kv_cache_usage_perc": 0.0,
      "prefix_cache_queries": 141069.0,
      "prefix_cache_hits": 92736.0,
      "gpu_utilization_avg": 86.0,
      "gpu_memory_used_mb_total": 113888.0,
      "gpu_memory_utilization_avg": 86.8896484375,

      "raw_notes": [ ]

    },
    "metrics_after": {
      "collected": true,
      "source": "vllm_metrics+nvidia_smi",
      "kv_cache_usage_perc": 0.0,
      "prefix_cache_queries": 162997.0,
      "prefix_cache_hits": 112432.0,
      "gpu_utilization_avg": 100.0,
      "gpu_memory_used_mb_total": 113888.0,
      "gpu_memory_utilization_avg": 86.8896484375,

      "raw_notes": [ ]

    }
  },
  {
    "model": "Qwen3-32B",
    "case_name": "step_15",
    "thinking_mode": "think",
    "context_tokens_target": 25037,
    "success": true,
    "http_status": 200,
    "error": "",
    "ttft_seconds": 15.394691379275173,
    "total_seconds": 24.415882788132876,
    "decode_seconds": 9.021191408857703,
    "raw_chunks": 129,
    "prompt_tokens_estimated": 25037,
    "completion_tokens_estimated": 182,
    "prompt_tokens_reported": null,
```

```
      "completion_tokens_reported": null,
      "total_tokens_reported": null,
      "output_tps_estimated": 20.174718809457733,
      "output_preview": "<think> 好的，我现在需要处理用户的问题，即检索链路在大促压测期间出现延迟抖动和局部召回空结果的情况。根据之前的工具调用和结果，我需要综合所有信息来确定根本原因并提出解决方案。 首先，回顾之前的步骤，用户已经通过多个工具调用来收...",
      "metrics_before": {
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
        "prefix_cache_queries": 162997.0,
        "prefix_cache_hits": 112432.0,
        "gpu_utilization_avg": 94.25,
        "gpu_memory_used_mb_total": 113888.0,
        "gpu_memory_utilization_avg": 86.8896484375,

        "raw_notes": [ ]

      },
      "metrics_after": {
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
        "prefix_cache_queries": 188083.0,
        "prefix_cache_hits": 134336.0,
        "gpu_utilization_avg": 95.0,
        "gpu_memory_used_mb_total": 113888.0,
        "gpu_memory_utilization_avg": 86.8896484375,

        "raw_notes": [ ]

      }
    },
    {
      "model": "Qwen3-32B",
      "case_name": "step_15",
      "thinking_mode": "no_think",
      "context_tokens_target": 25038,
      "success": true,
      "http_status": 200,
```

```
      "error": "",
      "ttft_seconds": 15.435323198325932,
      "total_seconds": 15.713504565414041,
      "decode_seconds": 0.27818136708810925,
      "raw_chunks": 7,
      "prompt_tokens_estimated": 25038,
      "completion_tokens_estimated": 4,
      "prompt_tokens_reported": null,
      "completion_tokens_reported": null,
      "total_tokens_reported": null,
      "output_tps_estimated": 14.379108284175869,
      "output_preview": "<think> </think> OK",
      "metrics_before": {
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
        "prefix_cache_queries": 188083.0,
        "prefix_cache_hits": 134336.0,
        "gpu_utilization_avg": 85.75,
        "gpu_memory_used_mb_total": 113888.0,
        "gpu_memory_utilization_avg": 86.8896484375,

        "raw_notes": [ ]

      },
      "metrics_after": {
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
        "prefix_cache_queries": 213173.0,
        "prefix_cache_hits": 156240.0,
        "gpu_utilization_avg": 99.25,
        "gpu_memory_used_mb_total": 113888.0,
        "gpu_memory_utilization_avg": 86.8896484375,

        "raw_notes": [ ]

      }
    },
    {
      "model": "Qwen3-32B",
```

```
      "case_name": "step_18",
      "thinking_mode": "think",
      "context_tokens_target": 27892,
      "success": true,
      "http_status": 200,
      "error": "",
      "ttft_seconds": 14.491989102680236,
      "total_seconds": 23.517217736691236,
      "decode_seconds": 9.025228634011,
      "raw_chunks": 129,
      "prompt_tokens_estimated": 27892,
      "completion_tokens_estimated": 177,
      "prompt_tokens_reported": null,
      "completion_tokens_reported": null,
      "total_tokens_reported": null,
      "output_tps_estimated": 19.611691534659492,
      "output_preview": "<think> 好的，我现在需要处理用户的问题，即检索链路在大促压测期间出现延迟抖动和局部召回空结果的情况。用户希望我像一个ReAct运维代理一样逐步调查，优先检查运行手册、最近变更、监控、日志、配置、工单和集群事件，再决定下一步。 ...",
      "metrics_before": {
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
        "prefix_cache_queries": 213173.0,
        "prefix_cache_hits": 156240.0,
        "gpu_utilization_avg": 98.0,
        "gpu_memory_used_mb_total": 113888.0,
        "gpu_memory_utilization_avg": 86.8896484375,

        "raw_notes": [ ]

      },
      "metrics_after": {
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
        "prefix_cache_queries": 240945.0,
        "prefix_cache_hits": 181312.0,
        "gpu_utilization_avg": 80.75,
        "gpu_memory_used_mb_total": 113888.0,
```

```
      "gpu_memory_utilization_avg": 86.8896484375,

      "raw_notes": [ ]

    }
  },
  {
    "model": "Qwen3-32B",
    "case_name": "step_18",
    "thinking_mode": "no_think",
    "context_tokens_target": 27894,
    "success": true,
    "http_status": 200,
    "error": "",
    "ttft_seconds": 14.497181296814233,
    "total_seconds": 14.781645926181227,
    "decode_seconds": 0.2844646293669939,
    "raw_chunks": 7,
    "prompt_tokens_estimated": 27894,
    "completion_tokens_estimated": 4,
    "prompt_tokens_reported": null,
    "completion_tokens_reported": null,
    "total_tokens_reported": null,
    "output_tps_estimated": 14.061502158989034,
    "output_preview": "<think> </think> OK",
    "metrics_before": {
      "collected": true,
      "source": "vllm_metrics+nvidia_smi",
      "kv_cache_usage_perc": 0.0,
      "prefix_cache_queries": 240945.0,
      "prefix_cache_hits": 181312.0,
      "gpu_utilization_avg": 80.75,
      "gpu_memory_used_mb_total": 113888.0,
      "gpu_memory_utilization_avg": 86.8896484375,

      "raw_notes": [ ]

    },
    "metrics_after": {
      "collected": true,
      "source": "vllm_metrics+nvidia_smi",
```

```
        "kv_cache_usage_perc": 0.0,
        "prefix_cache_queries": 268721.0,
        "prefix_cache_hits": 206384.0,
        "gpu_utilization_avg": 100.0,
        "gpu_memory_used_mb_total": 113888.0,
        "gpu_memory_utilization_avg": 86.8896484375,

        "raw_notes": [ ]

      }
    },
    {
      "model": "Qwen3-32B",
      "case_name": "step_20",
      "thinking_mode": "think",
      "context_tokens_target": 29425,
      "success": true,
      "http_status": 200,
      "error": "",
      "ttft_seconds": 8.403404055163264,
      "total_seconds": 17.44319086894393,
      "decode_seconds": 9.039786813780665,
      "raw_chunks": 129,
      "prompt_tokens_estimated": 29425,
      "completion_tokens_estimated": 177,
      "prompt_tokens_reported": null,
      "completion_tokens_reported": null,
      "total_tokens_reported": null,
      "output_tps_estimated": 19.580107766498774,
      "output_preview": "<think> 好的，我现在需要处理用户的问题，即检索链路在大促压测期间出现延迟抖动和局部召回空结果的情况。用户希望我像一个ReAct运维代理一样逐步调查，优先检查运行手册、最近变更、监控、日志、配置、工单和集群事件，再决定下一步。 ...",
      "metrics_before": {
        "collected": true,
        "source": "vllm_metrics+nvidia_smi",
        "kv_cache_usage_perc": 0.0,
        "prefix_cache_queries": 268721.0,
        "prefix_cache_hits": 206384.0,
        "gpu_utilization_avg": 100.0,
        "gpu_memory_used_mb_total": 113888.0,
```

```
      "gpu_memory_utilization_avg": 86.8896484375,

      "raw_notes": [ ]

    },
    "metrics_after": {
      "collected": true,
      "source": "vllm_metrics+nvidia_smi",
      "kv_cache_usage_perc": 0.0,
      "prefix_cache_queries": 297907.0,
      "prefix_cache_hits": 234128.0,
      "gpu_utilization_avg": 90.5,
      "gpu_memory_used_mb_total": 113888.0,
      "gpu_memory_utilization_avg": 86.8896484375,

      "raw_notes": [ ]

    }
  },
  {
    "model": "Qwen3-32B",
    "case_name": "step_20",
    "thinking_mode": "no_think",
    "context_tokens_target": 29427,
    "success": true,
    "http_status": 200,
    "error": "",
    "ttft_seconds": 7.992037368007004,
    "total_seconds": 8.284638602286577,
    "decode_seconds": 0.29260123427957296,
    "raw_chunks": 7,
    "prompt_tokens_estimated": 29427,
    "completion_tokens_estimated": 4,
    "prompt_tokens_reported": null,
    "completion_tokens_reported": null,
    "total_tokens_reported": null,
    "output_tps_estimated": 13.670482319900614,
    "output_preview": "<think> </think> OK",
    "metrics_before": {
      "collected": true,
      "source": "vllm_metrics+nvidia_smi",
```

```
      "kv_cache_usage_perc": 0.0,
      "prefix_cache_queries": 297907.0,
      "prefix_cache_hits": 234128.0,
      "gpu_utilization_avg": 90.5,
      "gpu_memory_used_mb_total": 113888.0,
      "gpu_memory_utilization_avg": 86.8896484375,

      "raw_notes": [ ]

    },
    "metrics_after": {
      "collected": true,
      "source": "vllm_metrics+nvidia_smi",
      "kv_cache_usage_perc": 0.0,
      "prefix_cache_queries": 327097.0,
      "prefix_cache_hits": 261888.0
```